


Hybrid backend: per-step reward shaping (NOT policy-loss augmentation)



Implementation: `envs/backends/surrogate_backend.py` lines 343–350; `training/train_thermostatic.py` line 436